

```
In [ ]: import numpy as np

arr = np.array([-3.7987907838045455e-05, -3.539867007685948e-06, -6.407242018099311e-14, -3.511661427459041e-14, 8.77
-2.2871226512677367e-07, 4.960411854604274e-07, 6.2967408793589835e-12, -1.7319479184152442e-14])

arr
```

```
In [10]: # df[0][0], df[0].values[0]

for index in df.index:
    print(index)
```

```
0
1
2
3
4
5
6
```

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

```
In [2]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                          n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make\_classification.html)
```

```
In [3]: X.shape, y.shape
```

```
Out[3]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
In [3]: #please don't change random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

```
In [4]: # Standardizing the data.
scaler = StandardScaler()
x_train = scaler.fit_transform(X_train)
x_test = scaler.transform(X_test)
```

```
In [5]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[5]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

SGD classifier

```
In [6]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15, penalty='l2', tol=1e-3, verbose=0)
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html)
```

```
Out[6]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                    random_state=15, verbose=2)
```

```
In [7]: clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.05 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.06 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.07 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.08 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.09 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.10 seconds.
Convergence after 10 epochs took 0.10 seconds
```

```
Out[7]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                    random_state=15, verbose=2)
```

```
In [8]: clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

```
Out[8]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
                  0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
                  0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
        (1, 15),
        array([-0.8531383]))
```

This is formatted as code

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in [def initialize_weights\(\)](#))
- Create a loss function (Write your code in [def logloss\(\)](#))

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in [def gradient_dw\(\)](#))

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^{(t)})) - \frac{\lambda}{N} w^{(t)}$$

- Calculate the gradient of the intercept (write your code in [def gradient_db\(\)](#)) [check this](#)

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^{(t)})$$

- Update weights and intercept (check the equation number 32 in the above mentioned [pdf](#)):

$$w^{(t+1)} \leftarrow w^{(t)} + \alpha(dw^{(t)})$$

$$b^{(t+1)} \leftarrow b^{(t)} + \alpha(db^{(t)})$$

- calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
- And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
- append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [9]: def initialize_weights(dim):
        ''' In this function, we will initialize our weights and bias'''
        #initialize the weights to zeros array of (1,dim) dimensions
        #you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
        #initialize bias to zero
        w= np.zeros_like((dim))
        b= 0
        return w,b
```

```
In [10]: dim = X_train[0]
w,b = initialize_weights(dim)
print('w =',(w))
print('b =',str(b))

w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

```
In [11]: dim = X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
    assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
    return True
grader_weights(w,b)
```

Out[11]: True

Compute sigmoid

$$\text{sigmoid}(z) = 1/(1 + \exp(-z))$$

```
In [12]: def sigmoid(z):
''' In this function, we will return sigmoid of z'''
# compute sigmoid(z) and return

return (1/(1 + np.exp(-1*z)))
```

Grader function - 2

```
In [13]: def grader_sigmoid(z):
val=sigmoid(z)
assert(val==0.8807970779778823)
return True
grader_sigmoid(2)
```

Out[13]: True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

```
In [14]: def logloss(y_true,y_pred):
'''In this function, we will compute log loss '''
n=len(y_true)
loss = 0.0
for (y_true,y_pred) in zip(y_true,y_pred):
loss += ((y_true*np.log10(y_pred)) + ((1-y_true) * np.log10(1-y_pred)))
loss = -1 *(loss/n)
return loss
```

```
In [24]: import math
def logloss(y_true,y_pred):
'''In this function, we will compute log loss '''
loss=0
for true,pred in zip(y_true,y_pred):
loss+=true*math.log10(pred)+(1-true)*math.log10(1-pred)
return -loss/len(y_true)
```

Grader function - 3

```
In [15]: def grader_logloss(true,pred):
    loss = logloss(true,pred)
    assert(loss==0.07644900402910389)
    return True
true=[1,1,0,1,0]
pred=[0.9,0.8,0.1,0.8,0.2]
grader_logloss(true,pred)
```

Out[15]: True

Compute gradient w.r.to 'w'

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

```
In [16]: # calculating gradient for single data point.
def gradient_dw(x,y,w,b,alpha,N):
    '''In this function, we will compute the gardient w.r.to w '''
    z = np.dot(x,w.T) + b
    dw = x*(y-sigmoid(z)) - (alpha*w)/N
    return dw
```

Grader function - 4

```
In [17]: def grader_dw(x,y,w,b,alpha,N):
    grad_dw = gradient_dw(x,y,w,b,alpha,N)
    assert(np.sum(grad_dw) == 2.613689585)
    return True
grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
-2.81434437, -0.86771071, -0.04073287, 0.84827878, 1.99451725,
3.67152472, 0.01451875, 2.01062888, 0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N = len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

Out[17]: True

Compute gradient w.r.to 'b'

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

```
In [18]: def gradient_db(x,y,w,b):  
        '''In this function, we will compute gradient w.r.to b '''  
        z = np.dot(x,w.T)+b  
        db = y - sigmoid(z)  
        return db
```

Grader function - 5

```
In [19]: def grader_db(x,y,w,b):  
        grad_db=gradient_db(x,y,w,b)  
        assert(grad_db== -0.5)  
        return True  
        grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,  
                        -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,  
                        3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])  
        grad_y=0  
        grad_w,grad_b=initialize_weights(grad_x)  
        alpha=0.0001  
        N=len(X_train)  
        grader_db(grad_x,grad_y,grad_w,grad_b)
```

Out[19]: True

Predict_function:

```
In [22]: def pred(w,b,X):  
        N = len(X)  
        predicted_prob = []  
        for i in range(N):  
            z=np.dot(w,X[i])+b  
            predicted_prob.append(sigmoid(z))  
        return np.array(predicted_prob)
```

Implementing logistic regression

```
In [59]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):  
        ''' In this function, we will implement logistic regression'''  
        #initialize the weights:
```



```

w,b = initialize_weights(X_train[0])
N = len(X_train)

# define lists to store train & test loss:
train_loss_list,test_loss_list = [],[]

for e in range(epochs):
    grad_w = 0
    grad_b = 0

    for row in range(N):# pass each instance of training data to update weights

        #compute gradient:
        grad_w = gradient_dw(X_train[row],y_train[row],w,b,alpha,N)
        grad_b = gradient_db(X_train[row],y_train[row],w,b)

        #update w & b:
        w = w + (eta0*grad_w)
        b = b + (eta0*grad_b)

    # using updated weights(each epoch) predict for X_train &X_test:
    pred_train = pred(w,b,X_train)
    pred_test = pred(w,b,X_test)

    # compute loss between predicted values & actual values:
    train_loss = logloss(y_train,pred_train)
    test_loss = logloss(y_test,pred_test)

    # append the loss obtained in each epoch:
    train_loss_list.append(train_loss)
    test_loss_list.append(test_loss)

return w,b,train_loss_list,test_loss_list

```

Observe Final weights :

```

In [68]: # call the train function:(defined above)
alpha= 0.0001
eta0= 0.0001
epochs= 10

```

```
w,b,trainloss,testloss = train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)

#Slope coefficients:(weights)
print(w)

# intercept value:
print(b)
```

```
[-0.42320236  0.19097504 -0.14588903  0.33813461 -0.21204107  0.56528021
 -0.44537758 -0.09169276  0.21798654  0.16980147  0.19524869  0.00226123
 -0.0778474   0.33881857  0.02215503]
-0.850591279771658
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```
In [69]: # these are the results we got after we implemented sgd and found the optimal weights and intercept
w=clf.coef_, b=clf.intercept_
```

```
Out[69]: (array([[ 0.00016455,  0.00549939,  0.00270133, -0.00330946, -0.00385437,
                  0.00511442,  0.00704724,  0.00239537,  0.00871335, -0.01103979,
                  -0.00180322, -0.00195793,  0.0017563 ,  0.00029055, -0.00051218]]),
          array([0.00254702]))
```

Observation : We can observe that the custom implementation & SGDClassifier's weights/intercept are as close enough with 10^{-3} difference.

Plot epoch number vs train , test loss


- epoch number on X-axis
- loss on Y-axis

```
In [99]: # Necessary libraries:
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('seaborn')

# Epoch number vs train,test loss:
epoch_number = [k+1 for k in range(epochs)]
sns.lineplot(x = epoch_number,y = trainloss,marker = "o")
```

```
sns.lineplot(x = epoch_number,y = testloss,marker = "o")
```

```
#customize plots:  
plt.xlabel("Epoch Number")  
plt.ylabel("Loss between Train & Test Set")  
plt.title("Epoch number vs Train-Test Loss")  
plt.legend(["Train loss","Test loss"])  
plt.show()
```



Observe the Train & Test accuracy:

```
In [98]: # Define a function to predict X_train & X_test with the final weights:  
def pred_accuracy(w,b,X):  
    N = len(X)  
    predict = []  
    for i in range(N):  
        z = np.dot(w,X[i])+b  
        if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
```

```
        predict.append(1)
    else:
        predict.append(0)
    return np.array(predict)

#observe the accuracy on train and test data:
print("The train accuracy :")
print(1-np.sum(y_train - pred_accuracy(w,b,X_train))/len(X_train))
print("The test accuracy :")
print(1-np.sum(y_test - pred_accuracy(w,b,X_test))/len(X_test))
```

```
The train accuracy :
0.9553333333333334
The test accuracy :
0.95288
```